



# Ensamble Learning

ITAM



# Outline

- Inspiration
  - Wisdom of the crowds
- Bagging
  - One instantiation: Random Forest
- Boosting
  - One instantiation: Adaboost

# Inspiration

- Galton's experiment
- Wisdom of crowds(Surowiecki)
- Who wants to be a millionaire

# Ensamble

- These techniques construct multiple models
  - Diverse models, diverse aspects of the data
- They combine the prediction of several models
  - Voting, average, etc.
- The differences between techniques boils down to differences in these two points

# Bagging

- Bagging (bootstrap aggregating)
- This method takes different random samples from the data (with replacement, bootstrap)
- Here we assume that different samples will yield to important differences in the resulting models
- The final prediction is formed using some combination rule
  - Voting (majority, plurality, weighed, ...)
  - Average (simple or weighed)

# Bagging

- Since diversity is key, another idea is to foster it by using different subsets of the attributes for each model (subspace sampling)
  - Increases diversity
  - Increases speed
  - Helps a bit with the curse of dimensionality
- Decision trees are sensitive to variations in the attributes so they make good candidates for this ideas to work
  - Small variation can lead to trees with different constructions
- The above is the basis for a technique called Random Forests

# Random Forest Training

- Input: Data  $D$ , number of trees  $T$ , number of attributes  $p$
- Output: A set of trees
- for  $t=1$  to  $T$  do
  - Create a sample  $D_t$  from  $D$  with replacement
  - Select  $p$  attributes at random and suppress the rest from  $D_t$
  - Grow a decision tree  $A_t$  using  $D_t$  without pruning
- return  $\{A_t \mid 1 \leq t \leq T\}$
- Note : It is recommended that  $p = \log(\text{number of attributes in } D)$  or  $p = \sqrt{p}$

# Random Forest Prediction

- Input: Set of trees, an instance  $x$  to label
- Output: Prediction for  $x$
- for  $t=1$  to  $T$  do
  - $y_t = A_t.\text{predict}(X)$
- if classification
  - return  $\text{vote}(\{y_t \mid 1 \leq t \leq T\})$
- else
  - return  $\text{mean}(\{y_t \mid 1 \leq t \leq T\})$



# Boosting

- This technique comes from the question of whether the classes of problems weakly learnable are strongly learnable are equivalent
  - The answer was given by Schapire in the article “The Strength of Weak Learnability”
  - The answer is Yes and the proof is by construction. The construction is called boosting (bootstrap aggregation)
- General idea:
  - Generate a set of models sequentially. Each new model is trained to correct the errors of the previous models. The output is a combination of them all

# AdaBoost

- Adaboost (adaptive boosting) is an implementation of this idea. In particular it establishes:
  - How to weigh the training examples to reflect the errors of the other models
  - How to weigh each model so as to reflect its importance and role in the final ensemble
- It uses an exponential cost function from which said weights are derived
- It takes as input any binary classification algorithm (assumes labels are -1 and 1)

# Boosting

- New weights for the training examples are computed in each iteration
  - This is implemented by sampling the training data with a new distribution
- The basic idea is to adjudicate half of the weight to the data that is correctly classified and the other half to the misclassified instances
  - Given the data  $D$ , each instance has an initial weight of  $1/|D|$
  - Subsequently given the classification error  $\epsilon$  we assign half of the weight to the correctly classified instances and half to the misclassified ones
    - For example if we are wrong on 25% of the instances, the next model will assign these double their weight while the correctly classified will be reduced to 0.66

# Boosting

---

**Algorithm 11.3:** Boosting( $D, T, \mathcal{A}$ ) – train an ensemble of binary classifiers from reweighted training sets.

---

**Input** : data set  $D$ ; ensemble size  $T$ ; learning algorithm  $\mathcal{A}$ .

**Output** : weighted ensemble of models.

```
1  $w_{1i} \leftarrow 1/|D|$  for all  $x_i \in D$ ; // start with uniform weights
2 for  $t = 1$  to  $T$  do
3   run  $\mathcal{A}$  on  $D$  with weights  $w_{ti}$  to produce a model  $M_t$ ;
4   calculate weighted error  $\epsilon_t$ ;
5   if  $\epsilon_t \geq 1/2$  then
6     set  $T \leftarrow t - 1$  and break
7   end
8    $\alpha_t \leftarrow \frac{1}{2} \ln \frac{1-\epsilon_t}{\epsilon_t}$ ; // confidence for this model
9    $w_{(t+1)i} \leftarrow \frac{w_{ti}}{2\epsilon_t}$  for misclassified instances  $x_i \in D$ ; // increase weight
10   $w_{(t+1)j} \leftarrow \frac{w_{tj}}{2(1-\epsilon_t)}$  for correctly classified instances  $x_j \in D$ ; // decrease weight
11 end
12 return  $M(x) = \sum_{t=1}^T \alpha_t M_t(x)$ 
```

---

# AdaBoost

- The choice of weights and alfa is related to minimizing the exponential cost function
- We wish to minimize

$$error = \sum_{i=1}^N e^{-m_i}$$
$$m_i = y_i \sum_{t=1}^T \alpha_t M_t(x_i)$$

Here  $y_i$  is the real class (1 or -1) and  $M_t$  is the predicted class

# Other methods

- The principle ideas of ensemble learning are:
  - Have different models that capture different patterns in the data
  - Have a way to combine them
- Having said this we could think of having an ensemble of methods using diverse techniques (neural networks + SVC + knn, etc.) with the idea to provide diversity
- We could think of combining them using yet another model, for instance a logistic regression
  - This is known as stacking
  - Now the mixing model has extra parameters to adjust (learn). You need to take this into account during the learning process

# Exercise

- Download data from
  - <http://archive.ics.uci.edu/ml/>
    - I suggest Abalone
  - <http://archive.ics.uci.edu/ml/datasets/Abalone?pagewanted=all>
- Use Sklearn and compare a decision tree a random forest (adaboost is optional)