



Machine Learning

SVM




Outline

- Support Vector Machines (SVM)
 - Intuition
 - Important tricks




Past classes

- We have seen how the basic idea of linear regression
 - Can be used for classification
 - Perceptron
 - Logistic regression
 - Can be used to approximate non linear functions
 - Transforming the independent variables
 - Adding attributes
 - Connecting many components to form a network




General Idea


- We want to keep on using the idea of separating classes with hyperplanes
 - We know how to solve this optimally
 - Problem, of course, is that not all problems are linearly separable
- Idea: Maybe not linearly separable in their current space, but perhaps they will be in a higher dimension
 - Create additional dimensions so that the categories become linearly separable
 - Similar in spirit to what we did adding attributes




Original Data




Increase in Dimension and Support Vectors





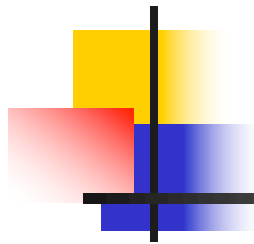
Optimal Margin Classifier





General Idea

- The idea is familiar, so are its problems
 - How do I add dimensions automatically? How many?
 - Adding dimensions is adding new attributes or features to the original data vectors
 - In having more dimensions we increment the sideeffects of the curse of dimensionality
 - More sensitivity to noisy data
 - Also, there is an increased risk of added computational burden
 - If the number of extra features is very large just computing them could be prohibitive



SVM


- Support vector machines are a proposal that addresses this issues
- Their development required the use of a few important ideas
 - We will review them presently



SVM


Main insights

- Optimal margin classifier
 - Reduces sensitivity to noise
 - Speeds up computations
 - Diminishes the curse of dimensionality
- Use of kernels
 - Adds dimensions to the data (arbitrarily large)
 - The kernel trick lets us pay a small computational price
- Soft margin
 - Slightly relax the requirement of lineal separability
 - Reduce sensitivity to noisy data




Optimal Margin Classifier


- Remember our lineal classifier?
 - In general, there are infinitely many hyperplanes to separate two classes
 - For what follows let us assume the classes are linearly separable in the current space
 - We will address this assumption later with the use of kernels and the soft margin



Binary Classificatio



¿Cuál separación es mejor, h1 h2 o h3?



Optimal Margin Classifier

- It is reasonable to assume that our prediction is more reliable the further away the data point is from the separating hyperplane
 - We are more confident that **A** belongs to the Star class than **B**
 - A small change in h_3 would change a change in classification for **B**
- h_2 generalizes better



Optimal Margin Classifier

Intuition

- If we think of our logistic classifier, the further away an instance is from the decision boundary, the closer the value of $g(w^T x + w_0)$ will be to 0 or 1 (depending on the class)
- If we think of the perceptron that has only two output values the same result is given to instances on the same side of the boundary
- The optimal margin combines both these ideas
 - We want instances within a band to satisfy the first observation
 - We want instances outside the band to satisfy the second; to not contribute to the error and permit the value of $w^T x + w_0$ to change without impact (as long as they remain on the proper side of the boundary)
 - In this way we can focus only on the instances that are on the border between classes



The Optimization Problem: the optimal margin classifier

- The cost function used will maximize this band
- The step transfer function establishes:
 - $V^{\wedge}(x) = 1$ if $w^T x \geq 0$ else -1
- What we now want
 - $V^{\wedge}(x) = 1$ if $w^T x \geq 1$ and
 - $V^{\wedge}(x) = -1$ if $w^T x \leq -1$
 - Otherwise we want the function to vary smoothly between -1 and 1



The Optimization Problem: the optimal margin classifier

$$\text{Min } \frac{1}{2} w^T w$$

tal que

$$y_i(w^T x_i + w_0) \geq 1 \quad \text{para } i = 1, \dots, M$$

With y_i 1 or -1




Geometric Interpretation


- Why does minimizing the size of w widen the margin?
 - The decision boundary is $w^T x^* = 0$ and is orthogonal to the weight vector w (x^* are the values of the features that satisfy the equation)
 - We require $w^T x^{(i)} \geq 1$ with $y=1$ and $w^T x^{(i)} \leq -1$ with $y=-1$
 - Remember that $w^T x^{(i)} = p||w|| = p \sqrt{\sum w_i^2}$ where p is the projection of vector $x^{(i)}$ onto vector w
 - We require $||w||^2$ to be small
 - The optimization problem forces the magnitude of the projection p to be large (very negative or very positive) in order to satisfy the restrictions thus pushing the decision boundary away from each $x^{(i)}$
 - For simplicity of explanation assume $w_0 = 0$ so that $||w||^2 = \sum w_j^2$




Geometric Interpretation



Which separates the classes better h1 or h2 given x?



Optimal Margin Classifier



The three circled vectors are the support vectors



Next Steps

- We want to find the vectors that characterize the optimal margin
 - This will help with the efficiency of the algorithm
 - This will help with generalization
- Increase the dimensionality of the problem so that the data points are linearly separable (or almost)
- For this...we need to transform our problem



Problem Transformation

- The dual problem
 - Transforms the original problem into another such that the solution to the new problem, under certain circumstances, is the solution to the original
- The goal is to have a problem that is easier or more efficient to solve
 - Our original optimization problem has an inequality that complicates the optimization (e.g. $|w^T x| \geq 1$)
 - It allows us to conserve only some vectors (the support vectors), the ones that characterize the optimal margin
 - It will allow the use of kernels for elevating the dimensionality of the problem in a computationally efficient manner



Problem Transformation

- We express the problem using Lagrange multipliers

$$L(w, w_0, \alpha) = \frac{1}{2} \|w\|^2 - \sum_{i=1}^m \alpha_i [y^{(i)} (w^T x^{(i)} + w_0) - 1]$$

- Note that the restrictions appear as a subtraction in the formula
- There is a set of conditions (KKT) under which the solution to the dual problem apply to the primal problem with the inequality restrictions
- This conditions give rise to the support vectors



Condiciones Karush-Kuhn-Tucker

1. $\frac{\partial}{\partial w_i} L(w, w_0, \alpha) = 0$
2. $\frac{\partial}{\partial w_0} L(w, w_0, \alpha) = 0$
3. $\alpha(y(w^T x + w_0) - 1) = 0$
4. $y(w^T x + w_0) - 1 \geq 0$
5. $\alpha \geq 0$



The Dual Optimization Problem

$$\max_{\alpha} W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j (x^{(i)})^T x^{(j)}$$

tal que $\alpha_i \geq 0, i = 1, \dots, m$

$$\sum_{i=1}^m \alpha_i y^{(i)} = 0$$

- To construct the optimization problem we minimize with respect to w by w_0 and substitute in the eqn for the lagrangian.
- If we solve the dual problem then we can recover the w 's (except w_0) using:

$$w = \sum_{i=1}^m \alpha_i y^{(i)} x^{(i)}$$

- This comes from differentiating and setting to zero (condition 1)
- w_0 is obtained similarly from condition 2



Support Vectors

- Note that we are still assuming data to be linearly separable
- Note that the problem maximizes α and forces it to be greater than or equal to zero
- Note that condition 3 forces many α to become zero. From here that only some vectors matter, those that support the separating hyperplane

$$\alpha(y(w^T x + w_0) - 1) = 0$$

- This $y(w^T x + w_0) - 1$ is 0 only for those vectors x that define the margin, for the rest α must be 0




The Dual Problem

- To make a prediction we need to compute $w^T x + w_0$. We substitute using the above equations

$$w^T x + w_0 = \left(\sum_{i=1}^m \alpha_i y^{(i)} x^{(i)} \right)^T x + w_0 = \sum_{i=1}^m \alpha_i y^{(i)} (x^{(i)})^T x + w_0$$

- This is our model
- If we have the α 's then the only expensive thing we need to do is compute the inner product of the support vectors with the input datum



Adding Dimensions

- This is relevant since there is a way to compute the inner product of vectors without having to manipulate them explicitly
- This, together with the fact that we will only use a few vectors (the support vectors) will allow us to increase the dimensionality at a low cost



Adding Dimensions

- Remember when we saw linear regression that we included features like x^2 in our data (we only had x as attribute)?
 - We talked about adding x^3 , $\sin(x)$ or whatever we could think of
 - One problem was that we could over fit
 - This will be mitigated since we only use a few vectors
 - Another that we did not have a principaled way to add them



Kernels

- What we are looking for is a way to increase the dimensionality of the data that allows us to compute the inner product efficiently
 - If we can do this we can add a lot of dimensions...even infinite
- This is what we will use the kernel trick for



Kernels

- The Kernel between two vectors is defined as:

$$K(x, z) = \phi(x)^T \phi(z)$$

- Where ϕ is a mapping of attributes to new features. For example, for a vector x in R

$$\phi(x) = \begin{bmatrix} x \\ x^2 \\ x^3 \end{bmatrix}$$

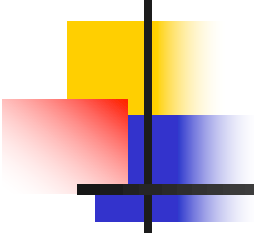


Kernels

Examples

- Suppose we have vectors x y z in R^3
- And suppose our mapping consists on multiplying every pair of components:

$$\phi(x) = \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_1 x_3 \\ x_2 x_1 \\ x_2 x_2 \\ x_2 x_3 \\ x_3 x_1 \\ x_3 x_2 \\ x_3 x_3 \end{bmatrix}$$



Kernels

Example

- The kernel

$$\begin{aligned} K(x, z) &= \phi(x)^T \phi(z) = \sum_{i=1}^n \sum_{j=1}^n x_i x_j z_i z_j \\ &= \left(\sum_{i=1}^n x_i z_i \right) \left(\sum_{j=1}^n x_j z_j \right) = (x^T z)^2 \end{aligned}$$

- Computing the kernel takes time proportional to the number of attributes while the number of generated features is proportional to the square of the number of attributes



Other Kernels

- Not all mappings have this property (Mercer's theorem establishes when a Kernel is valid).
- These are some common kernels

$$K(x, z) = (x^T z + c)^d$$

$$K(x, z) = \exp\left(\frac{-\|x - z\|^2}{2\sigma^2}\right)$$

$$K(x, z) = \tanh(\kappa x^T z - \partial)$$



Kernels

- Picking a Kernel and its parameters is a bit of an art
 - Not much theory
 - We could even design our own. Mercer's theorem will guide us to design a valid one
- One useful intuition to guide the selection of a Kernel is to think of $K(x,z)$ as a measure of how similar $\phi(x)$ and $\phi(z)$ are
 - The value of $K(x,z)$ will be large when they are similar and small when they are not



The New Model


$$\text{Prediccion}(x) = \sum_{i=1}^m \alpha_i y^{(i)} K(x^{(i)} x) + w_0$$

- Note that the model only uses the support vectors and that K only requires computing the inner product without enumerating the added attributes



Other Details

- All of the above was done assuming that the classes are linearly separable in some space
 - We can relax this supposition so that they are “almost” linearly separable and allow some datapoints to cross the border. This is called the soft margin
- There are a few efficient algorithms for solving the dual optimization problem
 - One of the most common is called “sequential minimal optimization” (SMO)



Soft Margin

- Data might not be linearly separable for two reasons:
 - The presence of noisy or mislabelled data
 - Data is intrinsically non-separable in this space
- For the second point we use Kernels to elevate the dimensionality
- For the first issue we use the soft margin



Soft Margin

- The idea of the soft margin is to permit some data points to be misclassified and not affect the margin
- We introduce a tolerance or slack parameter

- Before

$$y(w^T x + w_0) \geq 1$$

- Now

$$y(w^T x + w_0) \geq 1 - \xi$$

- We allow for data points to be “a little” misclassified, we allow their distance to the correct side of the boundary to be at most ψ



The Optimization Problem: soft margin classifier

$$\text{Min} \frac{1}{2} w^T w + C \sum_{i=1}^M \xi_i$$

tal que

$$y_i(w^T x_i + w_0) \geq 1 - \xi_i \quad \text{para } i = 1, \dots, M$$

$$\text{y } \xi_i \geq 0 \quad \text{para } i = 1, \dots, M$$




The Dual Optimization Problem

$$\max_{\alpha} W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j (x^{(i)})^T x^{(j)}$$

tal que $0 \leq \alpha_i \leq C, i = 1, \dots, m$

$$\sum_{i=1}^m \alpha_i y^{(i)} = 0$$

- Almost the same!
- The difference is that alfas must be smaller than or equal to a constant
- The support vector returned that are equal to C or $-C$ are those that don't respect the margin



Some advice from Andrew Ng

- If you have more attributes than data points
 - Use regularized logistic regression or SVMs with a linear kernel
- Few attributes and less than 10k examples
 - SVM with gaussian kernel
- Few attributes and more than 10k training examples
 - Add additional attributes and use logistic regression or SVM with linear kernel



Other comments

- The complexity of the resulting model is related to the relative number of support vectors (VC dimension)



Packages

- Implementing an SVM is a bit hard. A well tested and widely used library is libSVM



Exercise

- Download andSVM_2.csv
- Train a perceptron and plot the decision boundary
- Train an SVM using:
 - from sklearn.svm import SVC
 - Use a linear kernel
- Plot the data points and the decision boundary
- Plot the margin: the lines that pass through the support vectors
 - Try different values of C (1 and 100)



Exercise

- Create a data set in 2D x_1 y x_2 such that the points that lie within a circle are class 1 and those outside class 0
 - $x_1^2 + x_2^2 = r^2$
- Train a Neural net for this problem
 - Compute the confusion matrix and visualize the misclassified points
- Repeat using an SVM and compare